

1.

Message Passing:

- Explicit communication via messages
- Loose coupling of program components

Data Parallel:

- Address space is treated globally
- Set of tasks operate on data, but independently on disjoint partitions

Shared Memory:

- Implicit communication via memory operations
- Tight coupling of program components

Distributed Shared Memory:

- Assumes global memory address space that is logically partitioned
- Portions of shared memory may have an affinity for a particular process

2.

'Affinity' component in *upc_forall* function is $i/\text{THREADS}$. Iteration i will be performed by thread $[(i/\text{THREADS}) \% \text{THREADS}]$.

$i = 0, 1, 2, 3 \rightarrow$ thread 0
 $i = 4, 5, 6, 7 \rightarrow$ thread 1
 $i = 8, 9, 10, 11 \rightarrow$ thread 2
 $i = 12, 13, 14, 15 \rightarrow$ thread 3
 $i = 16, 17, 18, 19 \rightarrow$ thread 0
 And so on

X == N

Memory allocation:

Thread 0	Thread 1	Thread 2	Thread 3
image[0][0-63]	image[1][0-63]	image[2][0-63]	image[3][0-63]
image[4][0-63]	image[5][0-63]	image[6][0-63]	image[7][0-63]
...
image[60][0-63]	image[61][0-63]	image[62][0-63]	image[63][0-63]

Memory accesses:

	Thread 0	Thread 1	Thread 2	Thread 3
Shared	256	256	256	256
Remote	768	768	768	768

X == (N*N)/THREADS

Memory allocation:

Thread 0	Thread 1	Thread 2	Thread 3
image[0][0-63]	image[16][0-63]	image[0][0-63]	image[0][0-63]

image[1][0-63]	image[17][0-63]	image[1][0-63]	image[1][0-63]
...
image[15][0-63]	image[31][0-63]	image[47][0-63]	image[63][0-63]

Memory accesses:

	Thread 0	Thread 1	Thread 2	Thread 3
Shared	256	256	256	256
Remote	768	768	768	768

X == N/THREADS

Memory allocation:

Thread 0	Thread 1	Thread 2	Thread 3
image[0][0-15]	image[0][16-31]	image[0][32-47]	image[0][48-63]
image[1][0-15]	image[1][16-31]	image[1][32-47]	image[1][48-63]
...
image[63][0-15]	image[63][16-31]	image[63][32-47]	image[63][48-63]

Memory accesses:

	Thread 0	Thread 1	Thread 2	Thread 3
Shared	256	256	256	256
Remote	768	768	768	768

3.

```
#include <stdio.h>
#include <upc.h>
#define GNRL_SIZE 10//Assume the array size is 10
shared int arr0[GNRL_SIZE]; //shared array distributed across threads
shared int partial_sums[THREADS]; //shared array with one integer per thread
shared int sum; //sum will have affinity to thread 0
shared int mean; //mean will have affinity to thread 0
shared int var; //var will have affinity to thread 0

main(){
    int i, temp;
    upc_forall(i = 0; i < GNRL_SIZE; i ++; i){
        partial_sums[MYTHREAD] += arr0[i]; //calculate partial sums for each
thread
    }
    upc_barrier; //wait until all partial sums calculated

    if(MYTHREAD == 0){
        for(i = 0; i < THREADS; i ++){
            sum += partial_sums[i]; //sum the partial sums
        }
        mean = sum/GNRL_SIZE; //calculate mean
        for(i = 0; i < GNRL_SIZE; i ++){
            temp += (arr0[i]-mean)*(arr0[i]-mean);
        }
        var = temp/(GNRL_SIZE - 1); //calculate variance
        printf("Mean = %i\\tVariance = %i\\n", mean, var);
    }
}
```

4.

```
#include <stdio.h>
#include <upc_relaxed.h>

#define N 512 //assume image is 512x512
shared [N*N/THREADS] unsigned char img [N][N] //Shared array with block size
N*N/THREADS
shared int histogram[256];
upc_lock_t*lock;

void initialize(void){
    lock = upc_all_lock_alloc();
    CHECK_MEM(lock);
}

int main(void){
    int i, j;
    initialize();
    upc_barrier;
    upc_forall(i = 0; i < N; i ++; i*THREADS/N){
        for(j = 0; j < N; j ++){
            upc_lock(lock);
            histogram[img[i][j]] ++;
            upc_unlock(lock);
        }
    }
    upc_barrier;
    if(MYTHREAD == 0){
        //print histogram
    }
    return 0;
}
```